

Just Another Gibbs Sampler: JAGS

Inglis, A., Ahmed, A., Wundervald, B. and Prado, E.
PhD students at Hamilton Institute



14th November of 2018

Agenda

1. Introduction
2. R scripts
3. Python scripts
4. New models in JAGS
 - 4.1 Poisson model
 - 4.2 Exponential survival model
 - 4.3 Gaussian Mixture model
5. Final remarks
6. Acknowledgments

Introduction

- ▶ JAGS is a program to perform inference for Bayesian Hierarchical models. It was proposed by Martyn Plummer in 2003 as an alternative to the BUGS software;
- ▶ The Gibbs Sampler function of JAGS is the ARMS, which is flexible for dealing with univariate target densities;
- ▶ The main advantages of JAGS, compared to BUGS, is the programming language and its interfaces with other software, such as R and Python;
- ▶ Also, JAGS does not require the specification of the conditional distributions.

Introduction

- ▶ Although JAGS is a really useful software to learn, there are few good examples online that explain both theory and code;
- ▶ The goals of this project were:
 - ▶ Learn how to use JAGS to perform Bayesian modelling and;
 - ▶ Write new codes;
- ▶ Basically, JAGS requires only the sampling distribution and the prior distribution for each parameter.

Introduction

```
model {  
  # Likelihood  
  for (i in 1:n) {  
    y[i] ~ dnorm(mu[i], precision)  
    mu[i] <- alpha + beta * x  
  }  
  # Prior distributions  
  alpha ~ dnorm(0.0, 1.0E-3)  
  beta ~ dnorm(0.0, 1.0E-3)  
  aux <- dgamma(0.001, 0.001)  
  precision ~ 1.0/ aux  
}
```

R scripts

- ▶ Our main contributions were to add **mathematical details** and provide **real datasets examples** for **5** R scripts;
- ▶ The models involved were:
 - ▶ Random effect model;
 - ▶ Multivariate Normal model;
 - ▶ Beta regression;
 - ▶ Time series Beta Auto-Regressive model of order 1 and;
 - ▶ Mixture model.

R script - Beta regression

Let $\{Y_i\}_{i=1}^n$ be independent and identically distributed random variables and $X_i = (1, x_{i,1}, \dots, x_{i,1})$ a line vector with all covariates of the individual i . We assume that Y_i is distributed according to a Beta distribution, denoted by $Y_i \sim \text{Beta}(\mu, \phi)$, which may be written in the form

$$f(Y_i|\mu, \phi) = \frac{\Gamma(\phi)\Gamma(\mu\phi)}{\Gamma((1-\mu)\phi)} Y_i^{\mu\phi-1} (1 - Y_i)^{(1-\mu)/\phi},$$

where $0 < Y_i < 1$, $\mathbb{E}(Y_i) = \mu$, $\mathbb{V}(Y_i) = \mu(1-\mu)/(1+\phi)$, $0 < \mu < 1$ and $\phi > 0$. Thus, it is possible to model $g(\mu) = X_i\beta$, where $g(\cdot)$ is the link function that maps the unit interval into \mathbb{R} .

This parametrization was proposed by Ferrari and Cribari-Neto (2004).

R script - Beta regression

```
model
{
  # Likelihood
  for (t in 1:T) {
    y[t] ~ dbeta(a[t], b[t])
    a[t] <- mu[t] * phi
    b[t] <- (1 - mu[t]) * phi
    logit(mu[t]) <- alpha + beta * x[t]
  }
  # Priors
  alpha ~ dnorm(0, 10^-2)
  beta ~ dnorm(0, 10^-2)
  phi ~ dunif(0, 10)
}
```


R script - Beta regression

```
library(datasets)
head(attenu)


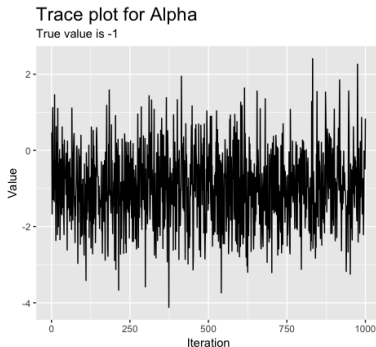
#Set up the data
acc=with(attenu,list(y=attenu$accel
                     ,T=nrow(attenu)))

# Set up jags model
jags_model=jags(acc,
                parameters.to.save = model_parameters,
                model.file = textConnection(model_code),
                n.chains=4,
                n.iter=1000,
                n.burnin=200,
                n.thin=2)

# Plot the jags output
print(jags_model)
```

R script - Beta regression

Simulation results for Beta regression, where three parameters were consider. The true values were set in $\alpha = -1$, $\beta = 0.2$ and $\phi = 5$.



Python scripts

- ▶ 3 Python scripts were created providing **mathematical details** and **simulated and real datasets examples**;
- ▶ The models involved were:
 - ▶ Bayesian linear regression;
 - ▶ Logistic regression;
 - ▶ Beta regression.

Python script - Logistic regression

The logistic regression model assumes that a sequence of independent and identically distributed random variables $\{Y_i\}_1^n$ has a Binomial distribution, denoted by $Y_i \sim \text{Binomial}(p_i)$, in the form of

$$f(Y_i|p_i) = \binom{n}{Y_i} p_i^{Y_i} (1 - p_i)^{1-Y_i},$$

where $Y_i \in \{0, 1\}$, $\log(\frac{p_i}{1-p_i}) = X_i\beta$, $X_i = (1, x_{i,1}, \dots, x_{i,1})$ is the line vector of covariates associated to the individual i and β is the vector of unknown parameters.

Python script - Logistic regression

```
model
{
  # Likelihood
  for (t in 1:n) {
    y[t] ~ dbin(p[t], 1)
    logit(p[t]) <- beta_0 + beta_1 * x_1[t] +
                  beta_2 * x_2[t]
  }

  # Priors
  beta_0 ~ dnorm(0.0,0.01)
  beta_1 ~ dnorm(0.0,0.01)
  beta_2 ~ dnorm(0.0,0.01)
}
```

Python script - Logistic regression

The data obtained for this section was adapted from data used to model logistic regression of moth mortalities when exposed to identical doses of a particular insecticide;

- ▶ Response variable: denotes the number of deaths observed in each batch;
- ▶ X_1 it is defined as the sexcode (i.e male or female);
- ▶ X_2 it is defined as the dose administered to each moth.

Python script - Logistic regression

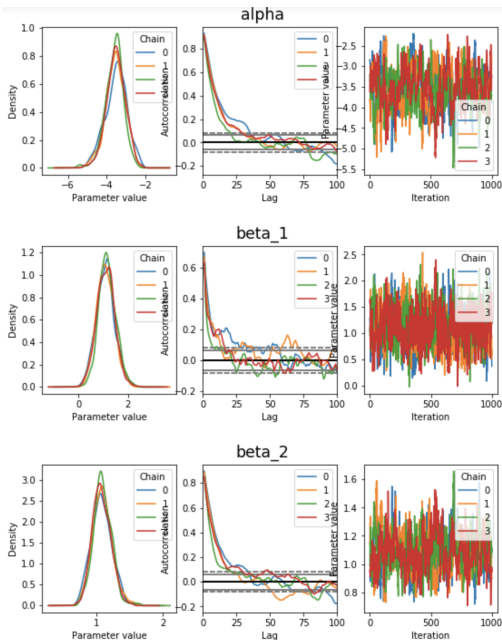
```
# Set up the data
model = pyjags.Model(code, data=dict(T = T, y = y,
                                     x_1 = x_1,
                                     x_2 = x_2, K = 1))

# Number of iterations to remove at start
model.sample(200, vars=[])

# Choose the parameters to watch and iterations:
samples = model.sample(1000, vars=['alpha', 'beta_1',
                                   'beta_2'])

print(jags_model)
```

Python script - Logistic regression



New models

Here we present 3 new models that we implemented both in R and Python;

The models are:

- ▶ Poisson regression model;
- ▶ Exponential survival model;
- ▶ Gaussian Mixture model.

Poisson regression model

A random variable Y is said to have a Poisson distribution with parameter λ if it takes integers $y = 0, 1, 2, \dots$ with probability mass function

$$P(Y = y) = \frac{\exp\{-\lambda\}\lambda^y}{y!},$$

where $\lambda > 0$. This mean can be modelled via a link function passed in a systematic component. For the Poisson regression case, the most widely used link function is the natural log, resulting in an equation that has the form

$$\log(\hat{\lambda}) = \beta_0 + \beta_1\phi(x_1) + \dots + \beta_n\phi(x_n).$$

Poisson regression model

```
model
{
  # Likelihood
  for (i in 1:T) {
    y[i] ~ dpois(p[i])
    log(p[i]) <- alpha + beta_1 * x_1[i] +
                  beta_2 * x_2[i]
  }
  # Priors
  alpha ~ dnorm(0.0, 0.01)
  beta_1 ~ dnorm(0.0, 0.01)
  beta_2 ~ dnorm(0.0, 0.01)
}
```

Poisson regression model

Simulate data -----

Some R code to simulate data from the Poisson model

```
T = 1000
```

```
set.seed(123)
```

```
x_1 = sort(runif(T, 0, 5))
```

```
x_2 = sort(runif(T, 0, 5))
```

```
alpha = 1
```

```
beta_1 = 1.2
```

```
beta_2 = -0.3
```



```
mu = alpha + beta_1 * x_1 + beta_2 * x_2
```

```
lambda = exp(mu)
```

```
y = rpois(n = T, lambda = lambda)
```

Poisson regression model

Y versus the explanatory variables with predicted line



Exponential Survival models

In survival analysis, we are usually interested in modelling the time until a certain event occurs. Let T be a random variable representing the survival times of individuals in some population.

$$F(t) = P(T \leq t) = \int_0^t f(u) du$$

Survival data is also often censored. In this case, the likelihood is written as

$$L(t) = \prod_{i=1}^n [f(t_i)]^{\delta_i} [S(t_i)]^{1-\delta_i},$$



where δ_i is the indicator variable that takes 1 for the failures and 0 for censored observations. We consider for the failure time the Exponential distribution, given by

$$f(t) = \frac{1}{\alpha} \exp\left\{-\frac{t}{\alpha}\right\}, \quad \alpha > 0.$$

Exponential Survival models

```
model
{
  # Likelihood
  for (i in 1:T) {
    mu[i] = exp(beta_1 * x_1[i] + beta_2 * x_2[i])
    t[i] ~ dexp(mu[i] * lambda_0)
  }
  # Priors
  lambda_0 ~ dgamma(1, 1)
  beta_1 ~ dnorm(0.0, 0.01)
  beta_2 ~ dnorm(0.0, 0.01)
}
```

Exponential Survival models



Gaussian Mixture model

The mixture model is viewed hierarchically: the observations y are modeled conditionally on the vector z , having z itself a probabilistic specification.

$$p(y_i|\theta, \lambda) = \lambda_1 f(y_i|\theta_1) + \cdots + \lambda_H f(y_i|\theta_H),$$


where the vector $\lambda = (\lambda_1, \dots, \lambda_H)$ represents the proportions of the population taken as being drawn from each $f_h(y_i|\theta_h)$ distribution, for $h = 1, \dots, H$, also that $\sum_{h=1}^H \lambda_h = 1$. Usually, the mixture components are assumed to be part of the same parametric family, such as the Gaussian, but with different parameter vectors. The unobserved variables is written as

$$z_{ih} = \begin{cases} 1, & \text{of the } i\text{th unite is drawn from the } h\text{th component} \\ 0, & \text{otherwise} \end{cases}$$

Gaussian Mixture model

```
model {  
  # Likelihood:  
  for(i in 1:N) {  
    y[i] ~ dnorm(mu[i] , 1/sigma_inv)  
    mu[i] <- mu_clust[clust[i]]  
    clust[i] ~ dcat(lambda_clust[1:Nclust])  
  }  
  # Prior:  
  sigma_inv ~ dgamma( 0.01 ,0.01)  
  mu_clust[1] ~ dnorm(0, 10)  
  mu_clust[2] ~ dnorm(5, 10)  
  
  lambda_clust[1:Nclust] ~ ddirch(ones)  
}
```

Gaussian Mixture model



Final remarks

- ▶ This project we learned how to use JAGS to perform Bayesian analysis;
- ▶ We had some challenges: i) JAGS was new for whole the group; ii) Difficulties to run the package **pyjags**;
- ▶ Our contributions were to provide mathematical details and codes for existing and new models. In the end, we produced:
 - ▶ 8 R scripts (3 new);
 - ▶ 3 Python scripts (all of them new);
- ▶ As future work other models can be implemented, such as other GLMs, Geostatistical models, more complex Survival models and other GLMMs with/without longitudinal structure.

Acknowledgments

This work was supported by a Science Foundation Ireland Career Development Award grant number: 17/CDA/4695



References

- Denison, David G. T., Christopher C. Holmes, Bani K. Mallick, and Adrian F. M. Smith. 2002. *Bayesian Methods for Nonlinear Classification and Regression*. Wiley.
- Gelman, Andrew, John B. Carlin, Hal S. Stern, and Donald B. Rubin. 2004. *Bayesian Data Analysis*. 2nd ed. Chapman; Hall/CRC.
- Kalbfleisch, J. D., and R. L. Prentice. 2002. *The Statistical Analysis of Failure Time Data*. 2nd ed. John Wiley & Sons.
- Ntzoufras, Ioannis. 2008. *Bayesian Modeling Using WinBUGS*. doi:10.1002/9780470434567.
- Royle, J Andrew, and Robert M Dorazio. 2008. *Hierarchical Modeling and Inference in Ecology: The Analysis of Data from Populations, Metapopulations and Communities*. Elsevier.
- Ferrari, Silvia, and Francisco Cribari-Neto. 2004. "Beta Regression for Modelling Rates and Proportions." *Journal of Applied Statistics* 31 (7). Taylor & Francis: 799–815.